

When and why a bug is introduced

Jon Eyolfson

University of Waterloo
jon@eyolfson.com

Abdel Maguid Tawakol

University of Waterloo
abdel.tawakol@gmail.com

Abstract

In this paper we attempt to find a correlation between the commit time of a code change and the rate of introduction of bugs. In other words, are developers more prone to bug-inducing changes during certain hours of the day or not. Before we start our search we were expecting our results to show that before lunch hour, in anticipation of going for lunch, and towards the end of the working day in anticipation of going home. Previous studies have explored whether there is a correlation between the bug-inducing changes rate and the day of the week, various mining techniques of bug-inducing changes, but to our knowledge no work has been to try and find a specific time slots during the day when developers might be more likely to introduce bugs. We mined data from Linux and Firefox using their software repositories. We developed an automated method for extracting information which finds a patch containing a fix and links it to patches which introduced the bug, which required the fix. We found the false positive rate to be under 20% after randomly sampling 100 reports. Our study has found that Thursday is a bad day to code, while Mondays are surprisingly good. Also between 12AM and 9AM, developers are more prone to introducing errors, while the least amount of bugs are introduced between 11AM and 3PM. Our data also shows that daily committers are less prone to introducing bugs, while day-job users are more prone. Finally, the bug lifetimes seem to decay exponentially, and on average is longer for larger projects.

1. Introduction

The area of bug detection and resolution is currently a very active area of research in the software related disciplines. The process of detecting and resolving software bugs is an extremely costly process because of the manpower required to perform quality assurance on the developed software and having developers perform the fixes. It is estimated that 20% - 30% of the time spent on a software project is spent on testing and integration [12]. In addition to the cost of manpower, the time required to complete these tasks can be costly on the release date of the software. As such, for the majority of software companies, the process of quality assurance is costly and time consuming. To put the cost into perspective, a study conducted by the National Institute of Standard and Technology in 2002 reported that software bugs cost the economy of the United States approximately \$59.5 billion annually [16]. The study also showed that approximately 64% of the costs of software bug fixes are incurred by the end users, making this issue of extreme importance for the consumers as well [16].

In our paper we will show the data collected from two large projects, which are complex and representative of most software projects, namely Linux and Firefox. Our data includes: percentage of bugs introduced vs total percentage of commits for every day of the week, percentage of bugs introduced every hour vs percentage of total commits every hour, percentage of introduction vs percentage of total commits by author classification, percentage of bugs introduced vs percentage of total commits grouped by author experience, and bug lifetime.

Our paper does not focus on solving a problem, but instead tries to find correlations between bug introductions and days of the week, hours of the day, commit introductions by author experience, commit introduc-

tions by author classification, and bug life time. This information can be used to try and reveal any potential correlations that could help leaders in the software industry determine ways to minimize any potential effects of working on a certain day, or during a certain hour. Collecting and analyzing data with regards to bug introduction will surely be beneficial to reducing bug introduction rates by examining some of the potential causes.

This paper will be organized in the following fashion: Section 2 will give an insight to our idea and how we approached the problem, Section 3 discusses our implementation for automatically finding bug introduction points, Section 4 presents the results we found, Section 5 will discuss the related academic work which our work compliments, Section 6 explores possible future work, and finally Section 7 concludes our work.

2. Idea & Approach

At the start of our research, we collected data manually from the Linux kernel repository using *git*. We examined logs for bug fixes, and traced back the bug-inducing changes to a specific user and time. This preliminary data was to be used for verifying the correctness of our automatically collected data and also give us an idea of what kind of results to expect.

Our solution attempts to find as many factors as possible that could potentially contribute to the introduction of bugs. This includes considering commit times of the changes as well as the authors making those changes. The goal is to find any meaningful correlations, and deduce recommendations to reduce the negative effects on bug introduction in the code. We built a database with all of the information that we have collected in order to easily extract data. Other researchers would also be able to use our data and find more correlations.

The two software projects we chose to start investigating are the Linux kernel and Firefox browser. Some of the questions we try to answer are:

- Is there a correlation between the time of day a commit is made and the number of bugs introduced?
- Is there a correlation between the classification of the user and the number of bugs introduced by the user?

- Are less experienced developers more prone to introducing bugs? (Light users will be considered less experienced)
- Is there a correlation between the day of the week and the number of bugs introduced by the users?
- Is there an improvement in the average bug lifetime? How is it distributed?

We intend to answer the questions posed above by looking at the code fixes and finding which commit caused the bug. The commits are gathered from the software repository, which in the case of Linux is *git* and for Firefox we converted the mercurial repository to *git* for ease of implementation. These contain the author and the time of the day the change was made, their time zone offset, and of course the actual change. We added this information to the database for future reference. The use of a database should be beneficial in the future for others to write queries and find data of interest to them, without having to go through the hassle of manually examining the software's repository. We also went through all of the commits and added more entries to the database which record the commits by each user. We then used this information to classify users into one of the following classifications: daily, weekly, monthly, or single committer.

We started off by manually reading the bug fixes and tracing them back to their source. After we have obtained enough experience doing this and have a collection of known bug introductions, we started to automate the process of extracting bug-inducing changes. We validated our data by randomly sampling automatically generated bug introductions and verifying them by hand to judge our technique.

After we constructed our database we wrote queries to answer the questions outlined previously. We then tried to explain some of these correlations and suggest some reasons as to why the bugs were introduced (this is detailed in the conclusion).

3. Implementation

We demonstrate our technique on a simple example, for instance say there is a change in the code repository such as the one in Listing 1. The listings contain a *unified diff* between the current version in the commit and the previous version, which is what we use in our tool. The bug here is that when *i* is equal to 256 the code

should actually use the *do_unicode* function. A commit which fixes this bug is shown in Listing 2.

Listing 1. An example bug introduction

```
Commit: f4ce718c ...
Message: I hope this works.
@@ -100,0 +100,5 @@
+     if (i <= 256) {
+         do_ascii(i);
+     else {
+         do_unicode(i);
+     }
```

Listing 2. An example bug fix

```
Commit: 2cdc03fe ...
Message: I fixed a bug!
@@ -100,1 +100,1 @@
-     if (i <= 256) {
+     if (i < 256) {
```

Our goal is to be able to find the commit which introduced the bug and also link it to the commit which included the fix. We found most developers indicate that their change is a fix by including the keyword "fix" in the commit message. We perform a simple keyword search with the word "fix" to find these commits. For this example we should find the commit starting with 2cdc03fe. From the unified diff we know line 100 was modified from the previous version on line 100. We then perform a *git-blame* on the previous version which indicates the last commit which changed the line. The output of the blame is shown in Listing 3

Listing 3. Blame of the previous version

```
f4ce718c ... 100     if (i <= 256) {
```

Using the blame information we found the original commit which introduced the bug, and can enter it into our database. This technique of blaming can be used for any modifications or removals easily, however additions pose a problem as the previous file gives us no information. In the case of additions we perform a blame on the current version of the file, we use the commit which modified the line before the new block of code as the commit which introduced the bug. In most cases the fix adds error checking code the original author forgot about.

We also take precautions when we analyze the changes. First we only look at changes which happen to C/C++ code. We ignore all comments as well since we are concerned solely about code changes. Finally, we

ignore whitespace when performing the blaming so we do not create a false introduction point from a change that didn't actually change the code.

Our process is automated using Python scripts which we wrote. We used bindings for git along with a diff library to interact with the code repository. We populate a MySQL database with the information we collect. A high level version of our database schema is shown in Listing 4.

Listing 4. Core database schema

```
class Repository:
    name = CharField()
    description = TextField()

class Author:
    repository =
        ForeignKey(Repository)
    name = CharField()
    email = CharField()

class Commit:
    author =
        models.ForeignKey(Author)
    sha1 = models.CharField()
    utc_time = models.DateTimeField()
    local_time =
        models.DateTimeField()

class Bug:
    introductions =
        ManyToManyField(Commit)
    fixes = ManyToManyField(Commit)
```

We also provided an extension to the database as an example of what can be done with the core data. Our extension for author information is shown in Listing 5. To classify each author we can perform the following steps: sort their commits by time, look at the time difference between consecutive commits (ignoring changes less than 30 minutes apart) and record the time difference as daily/weekly/monthly. We then classify them based on what time difference is between the majority of their commits. We also have an additional check for daily committers, if 85% of their commits are between 8 AM - 4 PM Monday to Friday we also classify them as having a day job on the project. Finally we determine how many months of experience they have

Type of Change	Amount (%)	
	Linux	Firefox
None	7.7 \pm 1.65	43.0 \pm 3.07
Add	11.0 \pm 1.94	6.9 \pm 1.57
Remove	2.0 \pm 0.87	1.8 \pm 0.82
Modify	35.7 \pm 2.97	13.8 \pm 2.14
Add/Modify	14.9 \pm 2.20	11.6 \pm 1.98
Add/Remove	5.8 \pm 1.45	2.3 \pm 0.93
Modify/Remove	5.9 \pm 1.46	3.2 \pm 1.09
All	17.0 \pm 2.33	17.4 \pm 2.35

Table 1. The types of changes between bug introductions and fixes from random sampling

by looking at the time difference from their first commit to last commit.

Listing 5. Extension to database schema

```
class AuthorInformation:
    CLASSIFICATIONS = (
        ('D', 'Daily'),
        ('W', 'Weekly'),
        ('M', 'Monthly'),
        ('S', 'Single'),
    )
    author = ForeignKey(Author)
    classification = CharField(
        choices=CLASSIFICATIONS)
    day_job = BooleanField()
    experience = IntegerField()
```

4. Results

We began by randomly sampling 1000 bugs from both Linux and Firefox and determined the types of changes, the results are shown in Table 1. Overall the percentages are similar for both software projects except for when there are no changes and changes which are only modifications. We believe this is because most of the bugs only applied to JavaScript for Firefox which we did not analyze, and they contain the class of bugs which are fixed by simple modifications.

We also noted that approximately 7.7% of the bug fixes for Linux were purely for comments, indicating that developers spend a non-trivial amount of time on comments. For Linux approximately 51% of the changes do not include additions. We manually checked the 49% more difficult cases to evaluate the effectiveness of our technique.

For Linux we manually checked a random sample of 50 bugs with the following types: 11 additions, 16 additions and modifications, 6 additions and removals, and 17 with all types of changes. We found 10 of the 50 to be false positives. Similarly for Firefox we randomly sampled 50 bugs with the following types: 9 additions, 3 additions and modifications, 18 additions and removals, and 20 with all types of changes. We found 13 of the 50 to be false positives.

There are five main sources of false positives we found, first we cover the rare cases. For one, the commit message we determined as a fix actually referred to fixing a merge conflict. Next, the blamed author was not correct due to them being the original author of the code, not the one who missed a check added at a later time. Finally a removal of dead code would trigger a non-existent bug. The more common cases involved reverting a change so it could be added in a later version and a code clean-up which involved moving functions or renaming. However in some cases you could argue that some of the code clean-ups should be counted as a bug.

Our false positive rates are currently between 20%-26% for the difficult cases of bug fixes. Since these account for approximately half of the overall fixes, we believe this is an upper bound on our false positive rate. Taking this into account we believe our data is representative and the results that we obtained are valid and meaningful.

First we determine if the day of the week has any impact on the likelihood of producing bugs. For the following graphs the percentage of introduction commits indicates the amount of commits out of the total number of commits which are bug introductions. Our results from Linux are shown in Figure 1. We see that Saturday is the worst day followed by Thursday, while Monday has the fewest percentage of bug introductions. The results for Firefox are shown in Figure 2 and agree with Thursday being one of the worse days while Monday is the best for committing changes which do not introduce bugs. A possible explanation is that developers rest over the weekend and have ample time to think about the problem before writing it on Monday when they know exactly what to do. It may also be correlated to the size of the change which we plan to investigate later as an addition to our technique.

Next we determine if the hour has any impact. Our results from Linux and Firefox are shown in Figure 3

and Figure 4 respectfully. Both show a significant increase in the amount of commits which introduce a bug between midnight and 9 AM. After this time there is a significant reduction in the amount of bug introductions between 11 AM and 3 PM. After 3 PM however, the likelihood of bug introductions fluctuates from hour to hour. This result is not very surprising since tired programmers are more likely to produce mistakes. It does however indicate the hours that programmers are at the peak of their productiveness in terms of not creating additional bugs.

We then investigated how each classification of developers fair in terms of their likelihood of creating a bug. The results from Linux and Firefox are shown in Figure 5 and Figure 6 respectfully. We see that developers who commit changes daily are a very significant portion of both projects and not developers that work on code as part of their day job. This trend is likely to be common to most open source projects. We also see for both projects the daily developers are less likely to produce bugs. On the other hand developers with this as their day job are more likely to produce bugs. The cause of this could be that the developers with this as their day job are required to make changes, while the daily developers are motivated purely by interest.

The final head-to-head comparison we performed is based on the amount of experience per author. We divided them up into 6 months intervals. Our results from Linux and Firefox are shown in Figure 7 and Figure 8 respectfully. For Linux authors with less than 2 years of experience are more likely to commit a bug while after 2 years the likelihood decreases. However the authors which committed throughout the history of the project are more likely to commit a bug which may be because they wrote the majority of the code. For Firefox, again the developers with less than 6 months of experience are more likely to commit a bug. There is a spike between 21 and 26 months as well. We believe these results may not be accurate due to the possibility of our experience metric being crude and not representative of the actual experience with the project.

Finally we determined the bug lifetimes for both projects to contrast them to previous work. We plotted the number of bugs for each lifetime divided into 4 month intervals. The results for Linux are shown in Figure 9, the average bug lifetime is 1.39 years. The results for Firefox are shown in Figure 10, with an average bug lifetime of 0.97 years. This shows a reduction in the av-




Figure 1. Linux percentage of bug introductions and percentage of total commits per day

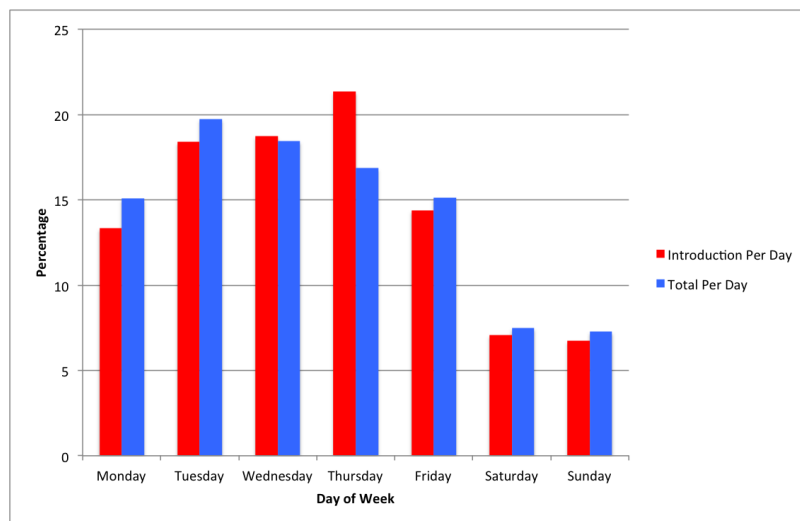


Figure 2. Firefox percentage of bug introductions and percentage of total commits per day




Figure 3. Linux percentage of bug introductions and percentage of total commits per hour

erage bug lifetime for Linux (in 2002 the average lifetime was 1.8 years). This indicates that developers are improving as their development process becomes more refined. We also see the average lifetime is lower for Firefox, this may be due to the complexity of the software, size of the software or the amount of users (the more users and complex, the more likely bugs will be revealed).

5. Related Work

No single solution has been found for solving the problem of bug introduction, but a lot of work has been done to try and minimize the problem, because it is unlikely that it could be completely resolved. Some of the things that could have been done is spend more time on software design to minimize the introduction of design related bugs by taking into account all the accumulated experiences to avoid repeating any mistakes that lead to bugs. Finally, companies have made strides in improving testing methodologies and automating the testing process as much as possible, which lead to savings of approximately \$22.5 billion [12]. The works presented below are mostly geared towards data collection for analysis, and trying to find potential sources that could

be contributing to bug-introducing changes to see what can be done to minimize bug-introduction.

When do changed induce fixes? One of the relevant academic works is this paper, which analyzes CVS achieves for fix-inducing changes. In other words, they examine code changes that lead to problems. They discuss a methodology to automatically locate fix-inducing changes by linking a version archive to a bug database such as BUGZILLA [15]. The authors examined the history for Mozilla and Eclipse for their data collection. The results they collected yielded that fix-inducing changes show distinct patterns with respect to their size and the day of the week they occurred. They discuss the general idea behind the process of finding fix-inducing changes as: 1. Start with a bug report in the bug database, indicating a fixed problem, extract the associated change from the version archive to get the location of the fix, and finally determine the earlier change at this location that was applied before the bug was reported. The authors describe syntactic and semantic analysis for the first step in the their process of identifying potential fixes. Finally, through the use of diff and annotate commands, as well as cycling through different versions of the code, the authors are




Figure 4. Firefox percentage of bug introductions and percentage of total commits per hour




Figure 5. Linux percentage of bug introductions and percentage of total commits per author classification

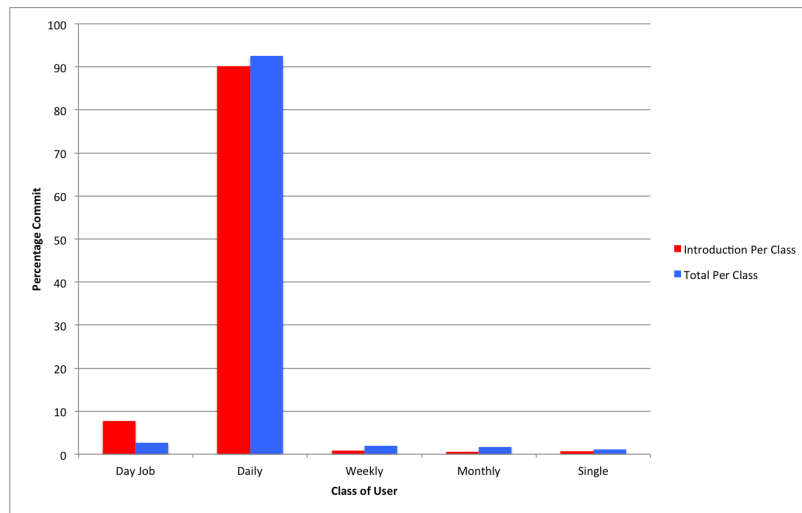


Figure 6. Firefox percentage of bug introductions and percentage of total commits per author classification




Figure 7. Linux percentage of bug introductions and percentage of total commits per author experience




Figure 8. Firefox percentage of bug introductions and percentage of total commits per author experience




Figure 9. Linux number of bugs against bug lifetimes in months




Figure 10. Firefox number of bugs against bug lifetimes in months

able to locate fix-inducing changes. The study concluded that most bugs are introduced on Fridays and Sundays. This work was important because it provided us with a guideline for a methodology for extracting fix-inducing changes.

Automatic identification of bug-introducing changes

The authors try to automate the process of finding bug-introducing changes, which would remove the manual work associated with going through bug reports or commit logs to collect this type of information. They use the SZZ algorithm, which traces from the location of the fix where the bug was introduced, and as such extract the time [7]. The weakness with this algorithm is that sometimes it makes false detections because not all modifications are fixes, and a moderate improvement from using just SZZ is the use of annotation graphs. Another improvement was ignoring format changes, which reduce a large number of false positives. This paper has inspired the approach we used for automating the process of identifying bug-introducing changes, which was key for collecting a large pool of data.

How long did it take to fix bugs? In this paper the author try to measure software quality as a function of the number of bugs. The authors examine the bug fix time of files in two open source software: ArgoUML and PostgreSQL, and tackle this issue by identifying when bugs are introduced and when the bugs are fixed [6]. The argument is files with the greatest bug-fix times, whose bug counts are greater than average, may need more attention to determine why bug fixes take such a long time potentially indicating the need for code refactoring to achieve faster bug fixes in the future. The authors first extracted change histories of the two projects, ArgoUML and PostgreSQL, using the Kenyon infrastructure. To identify a bug-fix, they searched for keywords such as fixed or bugs and they also searched for references to bug reports. They then applied the identified bug-introducing changes by applying the fix-inducing change identification algorithms described in the paper *When do changes induce fixes?* which I mentioned earlier. This work is of course relevant because it gives us a methodology for extracting bug-fix times.

If your bug database could talk This is another relevant work, where the authors perform experiments that

demonstrate how to relate developer, code, and process to defects in the code. This work tries to understand why some programs are more failure-prone than others. To answer this question, we have to know which programs are more failure-prone than others to search for properties of the programs or its development process that commonly correlate with defect density. The authors try to answer questions like can one predict failure-proneness from metrics like code complexity?, what does a high number of bugs found after release?, and do some developers write more failure-prone code than others?. After examining Eclipse database, some of the conclusions that the authors made were: new or combination of existing metrics need to be explored to study the relationship between complexity of code to the presence of bugs in a given class, it is difficult to predict post-release failures solely from process measurements, and there is a high variance in failure density in files owned by different developers [14]. Their methodology for extracting information from bug reports from BUGZILLA will be very useful for our project.

On the Nature of Commits This paper studies the nature of commits in two dimensions: define the size of commits in terms of number of files, and classify commits based on the content of their comments [3]. The authors investigated the distribution of commits according to the number of files, and their results show that the majority of commits contain a large number of files. The authors also developed a classification system for commits according to development and maintenance activities based on the content of their commits, a system that is more suitable for open source projects. Some of the major findings made by the authors include: the majority of the commits are not related to the development activities, corrective actions generate more tiny commits, and development activities are spread among all sizes of commits.

6. Future Work

There are a few parts of the implementation which could be improved to reduce the number of false positives. First, our fix detection for commit messages is very simplistic and could be improved to determine fixes which do not refer to any bugs. We can also introduce additional logic to ignore code which has been moved or renamed. The author information for experience also did not seem reliable, due to the fact that

we treat every name and e-mail pair as a unique author. However, this might not be the case if the authors simply changed their e-mail address.

We could also add additional extensions to the database including: categorizing authors by their official roles, and classifying the size of each commit as well to determine if larger commits are more likely to contain bugs.

In addition we would like to add additional software projects to the database. Another goal is to release the data to the community so others may extend and add to it. We believe there are much more interesting results which could be found using our database.

7. Conclusion

Resolving bugs represent a substantial amount of time and cost in software projects. It is important to investigate the cause of bugs in order to reduce the amount of bugs which occur. We believe that the data we have collected will be beneficial for the software quality assurance and software developing personnel to use the correlations we found, and investigate other possible correlations from the data we collected and stored in a database. From manually checking a random sample we found our data to be representative. Our data has revealed that Thursdays have a high bug introduction rate relative to the total number of commits for that day, while Mondays have the lowest bug introduction rates. There may be many reasons behind this correlation, but we believe it might be that people are being worn out towards the end of the week, and as such are more likely to introduce bugs, except this would mean that Fridays would be worst - which is not the case. It could also be that on Fridays people are more alert and focused because they are excited for the weekend. It was strange that Mondays yielded the best results for time to code, but this could be the case because people are just back from the weekend and they are fully charged, ready to work. Another possible theory for Monday's results, which is that people delegate easier tasks for Mondays because its the start of the week, or spend more time planning and laying out templates of what they will be coding, and as such are not as prone to introducing bugs. This of course could be proven by checking the size of the commits and what kind of changes are being made on Mondays. Our data also showed that coding between 12AM and 9AM is a bad idea, while the best time is between 11AM and

3PM. This is not a very surprising result, considering how most users sleep during these hours, and if they happen to be doing any coding during these hours, it would be considered to be outside of their usual working hours - making them more prone to errors. From our commits that were organized by author classification, it seems that daily committers are less prone to bug introduction than those who are classified as day job users. One theory that could potentially explain this result is that hobbyist working on open source projects do it because they are highly interested, while day job users are doing it to get paid, so the difference in motivations could negatively affect the performance of the developer. Finally our data shows that bug lifetimes decay exponentially, and on average are longer for larger projects. This result of course is positive because it indicates that the majority of bugs are dealt with fairly quickly.

References

- [1] C. Bird, A. Gourley, and P. Devanbu. Detecting patch submission and acceptance in oss projects. In *Proceedings of the Fourth International Workshop on Mining Software Repositories*, page 26. IEEE Computer Society, 2007. ISBN 076952950X.
- [2] G. Canfora and L. Cerulo. Impact analysis by mining software and change request repositories. 2005. ISSN 1530-1435.
- [3] L. Hattori and M. Lanza. On the nature of commits. In *Automated Software Engineering-Workshops, 2008. ASE Workshops 2008. 23rd IEEE/ACM International Conference on*, pages 63–71. IEEE, 2008.
- [4] J. Howison, M. Conklin, and K. Crowston. FLOSSmole: A collaborative repository for FLOSS research data and analyses. *International Journal of Information Technology and Web Engineering*, 1(3):17, 2006. ISSN 1554-1045.
- [5] H. Kagdi, M. Collard, and J. Maletic. A survey and taxonomy of approaches for mining software repositories in the context of software evolution. *Journal of Software Maintenance and Evolution: Research and Practice*, 19(2):77–131, 2007. ISSN 1532-0618.
- [6] S. Kim and E. Whitehead Jr. How long did it take to fix bugs? In *Proceedings of the 2006 international workshop on Mining software repositories*, pages 173–174. ACM, 2006.
- [7] S. Kim, T. Zimmermann, K. Pan, and E. Whitehead. Automatic identification of bug-introducing changes. In *21st IEEE/ACM International Conference on Automated Software Engineering, 2006. ASE’06*, pages 81–90, 2006.
- [8] S. Kim, T. Zimmermann, K. Pan, and E. Whitehead. Automatic identification of bug-introducing changes. In *21st IEEE/ACM International Conference on Automated Software Engineering, 2006. ASE’06*, pages 81–90, 2006.
- [9] Z. Li, L. Tan, X. Wang, S. Lu, Y. Zhou, and C. Zhai. Have things changed now?: an empirical study of bug characteristics in modern open source software. In *Proceedings of the 1st workshop on Architectural and system support for improving software dependability*, pages 25–33. ACM, 2006. ISBN 1595935762.
- [10] B. Livshits and T. Zimmermann. DynaMine: finding common error patterns by mining software revision histories. In *Proceedings of the 10th European software engineering conference held jointly with 13th ACM SIGSOFT international symposium on Foundations of software engineering*, pages 296–305. ACM, 2005. ISBN 1595930140.
- [11] L. Lopez-Fernandez, G. Robles, J. Gonzalez-Barahona, et al. Applying social network analysis to the information in cvs repositories. In *International Workshop on Mining Software Repositories*. Citeseer, 2004.
- [12] D. Reifer. Industry software cost, quality and productivity benchmarks. *DACS Newsletter*, 7(2), 2004.
- [13] K. Schneider, C. Gutwin, R. Penner, and D. Paquette. Mining a software developers local interaction history. In *MSR04: Proceedings of the 2004 international workshop on Mining software repositories*, pages 106–110. Citeseer, 2004.
- [14] A. Schröter, T. Zimmermann, R. Premraj, and A. Zeller. If your bug database could talk. In *Proceedings of the 5th international symposium on empirical software engineering*, volume 2, pages 18–20. Citeseer, 2006.
- [15] J. Sliwerski and T. Zeller. When do changes induce fixes? *Change*, 1(1.18):1–19, 2005.
- [16] G. Tassey. The economic impacts of inadequate infrastructure for software testing. *National Institute of Standards and Technology, RTI Project*, 2002.
- [17] A. Zaidman, B. Van Rompaey, S. Demeyer, and A. Van Deursen. Mining software repositories to study co-evolution of production & test code. In *Software Testing, Verification, and Validation, 2008 1st International Conference on*, pages 220–229. IEEE, 2008.